



## TCP and OSPF Assignment ELC 3080

Name	ID	Sec / BN
مجدي أحمد عباس عبد الحميد الابرق	9210899	3 / 36

Submitted To: Prof. Dr. Rania Osama  
3<sup>rd</sup> Year – Spring 2024

# Contents

1. TCP Lab.....	3
1.1. Effect of TCP Window Size .....	3
1.1.1. Vary iperf3 window size from 1Kbytes to 6 Kbytes in increments of 1 Kbyte, then set it to 12, 16, 24, 32 Kbytes. Plot the average throughput and the average number of retransmissions as function of window size. Comment on the results and explain the zigzag behavior noticed for larger window sizes. Note that number of retransmissions is the 5th column in iperf3 default output & throughput is the 4 <sup>th</sup> . You can plot data using MATLAB/python/any online graphing tool (Ex. <a href="https://chart-studio.plotly.com/create/#/">https://chart-studio.plotly.com/create/#/</a> ). The window size is the last option in the iperf command (return to INTRO Lab). .....	3
1.1.2. Click on any of the TCP a data segment whose source is node n7, dissect the segment by following different protocol layer headers from TCP->IP->Ethernet identifying how many header bytes are added by each layer, identify the TCP options used. ....	5
1.1.3. Repeat for an ACK packet sent from node n11.....	6
1.2. TCP short VS long paths .....	7
1.2.1. Compare the result of throughput with the case when connection was made to node n11. Why throughput drops when connecting to n8 although capacities on the two paths are the same. ....	7
1.3. Higher link Capacity with Drops VS Reliable Lower Capacity .....	7
1.3.1. Compare throughputs in cases a, b, c. Why b is better than c? .....	9
1.3.2. Compare throughputs in cases b, c and d. Which is better? Why? .....	9
1.3.3. Compare throughputs in cases e and f? Which is better? Why? .....	9
2. OSPF Lab.....	9
2.1. OSPF Link Cost Change.....	9
2.1.1. Check what happens to the path between n7 and n11 (as seen after steps 3 and 6)? Explain what happens. To help visualize the path choose toolbar->widgets-> Throughput.....	10
2.1.2. Set the cost of eth1 at node n5 back to 10. Establish two iperf3 connections: one from n7 to n11 and the second from n11 to n7 both for duration of 500 seconds. Now go to node n4 and set interface cost for interface connecting n4 with n5 to 40. What happens in the paths of the two connections? Explain what happens. What do you conclude? .....	11
2.2. OSPF Database Updates .....	12
2.2.1. Capture and explain the outputs due to execution of step 2. Why some destinations have more than route in the routing tables? .....	12
2.2.2.....After executing step 3, determine how long it took the network to exchange link state packets and adjust routing tables. (Hint: you can calculate the required time by observing the time of first OSPF update message and the last ACK from Wireshark). ....	13
2.2.3. After execution of step 4, identify the new routing table and router database at router n2. Explain the updates in the new routing table and the new database.....	13

# 1. TCP Lab

Each of the following sections describes an experiment based on the network topology provided by the [Project OSPF TCP.imn](#) file. In your report, divide the answers to correspond to one of these sections in exactly the same sequence and same name, for example use Question 2.1, Question 3.2 and so on. The [Introductory Lab](#) explained and introduced most of the needed tools like iperf3, vtysh, wireshark, etc. If you have not gone through it, do not start working on the assignment below. It is better to start the network emulation once and run Wireshark captures. You can reset Wireshark captures from experiment to another.

## 1.1. Effect of TCP Window Size

- Set wireshark filter to display TCP packets only.
- Start an iperf3 server on node n11.
- Start an iperf3 client on node n7 connecting to server on node n11 for a duration of 40 seconds and reporting interval of 10. **Note that in iperf, the client is the node sending the traffic and the server simply receives and sends an ACK.**

**1.1.1. Vary iperf3 window size from 1Kbytes to 6 Kbytes in increments of 1 Kbyte, then set it to 12, 16, 24, 32 Kbytes. Plot the average throughput and the average number of retransmissions as function of window size. Comment on the results and explain the zigzag behavior noticed for larger window sizes. Note that number of retransmissions is the 5th column in iperf3 default output & throughput is the 4th. You can plot data using MATLAB/python/any online graphing tool (Ex. <https://chart-studio.plotly.com/create/#/>). The window size is the last option in the iperf command (return to INTRO Lab).**




Figure 1: Throughput (Mbps) vs Window Size (kB)




Figure 2: Avg Number of Retransmissions vs Window Size (kB)

### Comment on your findings:

We observe from the throughput versus window size graph that there is a zigzag pattern as the window size increases from 1k to 5k. During this range, the number of retransmissions remains zero because the receiver buffer is not yet full. Consequently, as the window size increases, the throughput also rises until the receiver buffer becomes saturated at a window size of 5k. At the window size of 6k, the graph shows a sudden drop in throughput and an increase in retransmissions due to a significant number of packet losses. When this occurs, the transmitter identifies the issue and adjusts by allocating more space in the buffer. As a result, within a window size range between 6k and 12k, the throughput increases again along with a rise in retransmissions. However, if the window size exceeds 16k, the receiver buffer fills up once more, leading to packet losses and a subsequent decrease in throughput and decrease in number of retransmissions.

Table 1: Window Size Table Showing Throughput and Avg Number of Retransmissions Values

Figure 3: 1 kB Window Size

```

root@n7:/tmp/pycore.34469/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 1k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38557 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  1.34 MBytes  1.12 Mbits/sec  0    5.45 KBytes
[ 4] 10.00-20.00  sec  1.33 MBytes  1.12 Mbits/sec  0    5.45 KBytes
[ 4] 20.00-30.00  sec  1.33 MBytes  1.12 Mbits/sec  0    5.45 KBytes
[ 4] 30.00-40.00  sec  1.33 MBytes  1.12 Mbits/sec  0    5.45 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  5.33 MBytes  1.12 Mbits/sec  0
[ 4] 0.00-40.00  sec  5.33 MBytes  1.12 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34469/n7.conf#

```

Throughput = 1.12 Mbps, Retransmissions = 0

Figure 4: 2 kB Window Size

```

root@n7:/tmp/pycore.34472/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 2k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38560 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  3.46 MBytes  2.90 Mbits/sec  0    7.07 KBytes
[ 4] 10.00-20.00  sec  3.44 MBytes  2.89 Mbits/sec  0    7.07 KBytes
[ 4] 20.00-30.00  sec  3.46 MBytes  2.90 Mbits/sec  0    7.07 KBytes
[ 4] 30.00-40.00  sec  3.47 MBytes  2.91 Mbits/sec  0    7.07 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  13.8 MBytes  2.90 Mbits/sec  0
[ 4] 0.00-40.00  sec  13.8 MBytes  2.90 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34472/n7.conf#

```

Throughput = 2.9 Mbps, Retransmissions = 0

Figure 5: 3 kB Window Size

```

root@n7:/tmp/pycore.34475/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 3k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38563 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  6.86 MBytes  5.75 Mbits/sec  0   14.1 KBytes
[ 4] 10.00-20.00  sec  6.93 MBytes  5.81 Mbits/sec  0   14.1 KBytes
[ 4] 20.00-30.00  sec  6.93 MBytes  5.82 Mbits/sec  0   14.1 KBytes
[ 4] 30.00-40.00  sec  6.92 MBytes  5.81 Mbits/sec  0   14.1 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  27.6 MBytes  5.80 Mbits/sec  0
[ 4] 0.00-40.00  sec  27.6 MBytes  5.80 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34475/n7.conf#

```

Throughput = 5.8 Mbps, Retransmissions = 0

Figure 6: 4 kB Window Size

```

root@n7:/tmp/pycore.34478/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38566 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  6.93 MBytes  5.82 Mbits/sec  0   14.1 KBytes
[ 4] 10.00-20.00  sec  6.93 MBytes  5.81 Mbits/sec  0   14.1 KBytes
[ 4] 20.00-30.00  sec  6.92 MBytes  5.80 Mbits/sec  0   14.1 KBytes
[ 4] 30.00-40.00  sec  6.94 MBytes  5.82 Mbits/sec  0   14.1 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  27.7 MBytes  5.81 Mbits/sec  0
[ 4] 0.00-40.00  sec  27.7 MBytes  5.81 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34478/n7.conf#

```

Throughput = 5.81 Mbps, Retransmissions = 0

Figure 7: 5 kB Window Size

```

root@n7:/tmp/pycore.34481/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 5k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38569 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  6.94 MBytes  5.82 Mbits/sec  0   14.1 KBytes
[ 4] 10.00-20.00  sec  6.94 MBytes  5.82 Mbits/sec  0   14.1 KBytes
[ 4] 20.00-30.00  sec  6.94 MBytes  5.82 Mbits/sec  0   14.1 KBytes
[ 4] 30.00-40.00  sec  6.93 MBytes  5.82 Mbits/sec  0   14.1 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  27.7 MBytes  5.82 Mbits/sec  0
[ 4] 0.00-40.00  sec  27.7 MBytes  5.82 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34481/n7.conf#

```

Throughput = 5.82 Mbps, Retransmissions = 0

Figure 8: 6 kB Window Size

```

root@n7:/tmp/pycore.34487/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 6k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38575 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  164 KBytes  134 Kbits/sec  93  4.24 KBytes
[ 4] 10.00-20.00  sec  158 KBytes  130 Kbits/sec  95  4.24 KBytes
[ 4] 20.00-30.00  sec  158 KBytes  130 Kbits/sec  96  4.24 KBytes
[ 4] 30.00-40.00  sec  153 KBytes  125 Kbits/sec  94  4.24 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  634 KBytes  130 Kbits/sec  378
[ 4] 0.00-40.00  sec  624 KBytes  128 Kbits/sec
iperf Done.
root@n7:/tmp/pycore.34487/n7.conf#

```

Throughput = 0.13 Mbps, Retransmissions = 378

Figure 9: 12 kB Window Size

```

root@n7:/tmp/pycore.34493/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 12k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38581 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  1.97 MBytes  1.65 Mbits/sec  321  4.24 KBytes
[ 4] 10.00-20.00  sec  1.96 MBytes  1.65 Mbits/sec  316  4.24 KBytes
[ 4] 20.00-30.00  sec  1.96 MBytes  1.65 Mbits/sec  316  4.24 KBytes
[ 4] 30.00-40.00  sec  1.96 MBytes  1.65 Mbits/sec  316  4.24 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  7.86 MBytes  1.65 Mbits/sec  1269
[ 4] 0.00-40.00  sec  7.84 MBytes  1.64 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34493/n7.conf#

```

Throughput = 1.65 Mbps, Retransmissions = 1269

Figure 10: 16 kB Window Size

```

root@n7:/tmp/pycore.34499/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 16k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38587 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  1.99 MBytes  1.67 Mbits/sec  299  2.83 KBytes
[ 4] 10.00-20.00  sec  2.04 MBytes  1.71 Mbits/sec  297  4.24 KBytes
[ 4] 20.00-30.00  sec  2.04 MBytes  1.71 Mbits/sec  296  4.24 KBytes
[ 4] 30.00-40.00  sec  2.04 MBytes  1.71 Mbits/sec  296  4.24 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  8.12 MBytes  1.70 Mbits/sec  1188
[ 4] 0.00-40.00  sec  8.10 MBytes  1.70 Mbits/sec
iperf Done.
root@n7:/tmp/pycore.34499/n7.conf#

```

Throughput = 1.7 Mbps, Retransmissions = 1188

Figure 11: 24 kB Window Size

```

root@n7:/tmp/pycore.34502/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 24k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38590 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  329 KBytes  270 Kbits/sec  120  5.66 KBytes
[ 4] 10.00-20.00  sec  277 KBytes  227 Kbits/sec  114  8.48 KBytes
[ 4] 20.00-30.00  sec  297 KBytes  243 Kbits/sec  118  5.66 KBytes
[ 4] 30.00-40.00  sec  277 KBytes  227 Kbits/sec  113  7.07 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  1.15 MBytes  242 Kbits/sec  465
[ 4] 0.00-40.00  sec  1.11 MBytes  233 Kbits/sec
iperf Done.
root@n7:/tmp/pycore.34502/n7.conf#

```

Throughput = 0.242 Mbps, Retransmissions = 465

Figure 12: 32 kB Window Size

```

root@n7:/tmp/pycore.34505/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 32k
Connecting to host 10.0.13.20, port 5201
[ 4] local 10.0.12.20 port 38593 connected to 10.0.13.20 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr    Cwnd
[ 4] 0.00-10.00  sec  274 KBytes  225 Kbits/sec  111  7.07 KBytes
[ 4] 10.00-20.00  sec  161 KBytes  132 Kbits/sec  96  21.2 KBytes
[ 4] 20.00-30.00  sec  242 KBytes  198 Kbits/sec  119  7.07 KBytes
[ 4] 30.00-40.00  sec  161 KBytes  132 Kbits/sec  96  21.2 KBytes
-----
[ ID] Interval      Transfer    Bandwidth   Retr
[ 4] 0.00-40.00  sec  839 KBytes  172 Kbits/sec  422
[ 4] 0.00-40.00  sec  795 KBytes  163 Kbits/sec
iperf Done.
root@n7:/tmp/pycore.34505/n7.conf#

```

Throughput = 0.172 Mbps, Retransmissions = 422

### 1.1.2. Click on any of the TCP a data segment whose source is node n7, dissect the segment by following different protocol layer headers from TCP->IP->Ethernet identifying how many header bytes are added by each layer, identify the TCP options used.

798 65.274124 10.0.12.20 10.0.13.20 TCP 626 [TCP Fast Retransmission] 38599 > targus-...data1 [PSH, ACK] Seq=19568 Ack=1 Win=1116 Len=558 TSval=1137763 TSecr=1137762

Frame 798: 626 bytes on wire (5008 bits), 626 bytes captured (5008 bits)

Arrival Time: May 17, 2024 16:21:54.058059000 EEST  
Epoch Time: 1715952114.058059000 seconds  
[Time delta from previous captured frame: 0.000481000 seconds]  
[Time delta from previous displayed frame: 0.000481000 seconds]  
[Time since reference or first frame: 65.274124000 seconds]  
Frame Number: 798  
Frame Length: 626 bytes (5008 bits)  
Capture Length: 626 bytes (5008 bits)  
[Frame is marked: False]  
[Frame is ignored: False]  
[Protocols in frame: sll:ip:tcp:data]  
[Coloring Rule Name: Bad TCP]  
[Coloring Rule String: tcp.analysis.flags]

Linux cooked capture

Packet type: Sent by us (4)  
Link-layer address type: 1  
Link-layer address length: 6  
Source: 00:00:00:aa:00:18 (00:00:00:aa:00:18)  
Protocol: IP (0x0800)

Internet Protocol Version 4, Src: 10.0.12.20 (10.0.12.20), Dst: 10.0.13.20 (10.0.13.20)

Version: 4  
Header length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))  
0000 00.. = Differentiated Services Codepoint: Default (0x00)  
.... ..00 = Explicit Congestion Notification: Not-ECT (Not ECN-Capable Transport) (0x00)  
Total Length: 610  
Identification: 0x89e0 (35296)

Flags: 0x02 (Don't Fragment)  
0... .... = Reserved bit: Not set  
1... .... = Don't fragment: Set  
..0. .... = More fragments: Not set  
Fragment offset: 0  
Time to live: 64  
Protocol: TCP (6)

Header checksum: 0x818e [correct]  
[Good: True]  
[Bad: False]  
Source: 10.0.12.20 (10.0.12.20)  
Destination: 10.0.13.20 (10.0.13.20)

Transmission Control Protocol, Src Port: 38599 (38599), Dst Port: targus-getdata1 (5201), Seq: 19568, Ack: 1, Len: 558

Source port: 38599 (38599)  
Destination port: targus-getdata1 (5201)  
[Stream index: 3]  
Sequence number: 19568 (relative sequence number)  
[Next sequence number: 20126 (relative sequence number)]  
Acknowledgement number: 1 (relative ack number)  
Header length: 32 bytes

Flags: 0x018 (PSH, ACK)  
000. .... = Reserved: Not set  
...0 .... = Nonce: Not set  
.... 0... = Congestion Window Reduced (CWR): Not set  
.... .0.. = ECN-Echo: Not set  
.... ..0. = Urgent: Not set  
.... ...1 = Acknowledgement: Set  
.... .... 1... = Push: Set  
.... .... .0.. = Reset: Not set  
.... .... ..0. = Syn: Not set  
.... .... ...0 = Fin: Not set  
Window size value: 1116  
[Calculated window size: 1116]  
[Window size scaling factor: 1]

Checksum: 0x2f7c [validation disabled]  
[Good Checksum: False]  
[Bad Checksum: False]

Options: (12 bytes)  
No-Operation (NOP)  
No-Operation (NOP)

Timestamps: TSval 1137763, TSecr 1137762  
Kind: Timestamp (8)  
Length: 10  
Timestamp value: 1137763  
Timestamp echo reply: 1137762

SEQ/ACK analysis  
[Bytes in flight: 558]

TCP Analysis Flags  
[This frame is a (suspected) fast retransmission]  
[Expert Info (Warn/Sequence): Fast retransmission (suspected)]  
[Message: Fast retransmission (suspected)]  
[Severity level: Warn]  
[Group: Sequence]  
[This frame is a (suspected) retransmission]

Note:

The timestamp option includes a timestamp from the sender, which is then echoed by the receiver in each packet once it's enabled during connection setup. This helps calculate round-trip time samples to estimate packet loss and extends the 32-bit sequence number. On fast connections, sequence numbers can reset quickly, causing potential confusion between old and new data. The PAWS (Protection Against Wrapped Sequence numbers) system resolves this by discarding segments with outdated timestamps.

Figure 13: TCP Packet

Comment on your findings here:

- Frame Length: → 626 bytes.
- IP Header: → 20 bytes
- TCP Header: → 32 bytes (12 Option + 20 Without Option)
- TCP Options: → 12 bytes (2 NOP + 10 Timestamp)
- TCP Options Timestamp: → 10 bytes (4T<sub>x</sub> + 4R<sub>x</sub> + 1 Length + 1 Kind)
- Ethernet Header: = Frame Length – IP Header – TCP Header = 626 – 20 – 32 = 574 bytes
- As Shown in Fig. 13: Packet Sent from n7 to n11 has large Size (626 B) as it contains Data & Headers.



### 1.1.3. Repeat for an ACK packet sent from node n11.

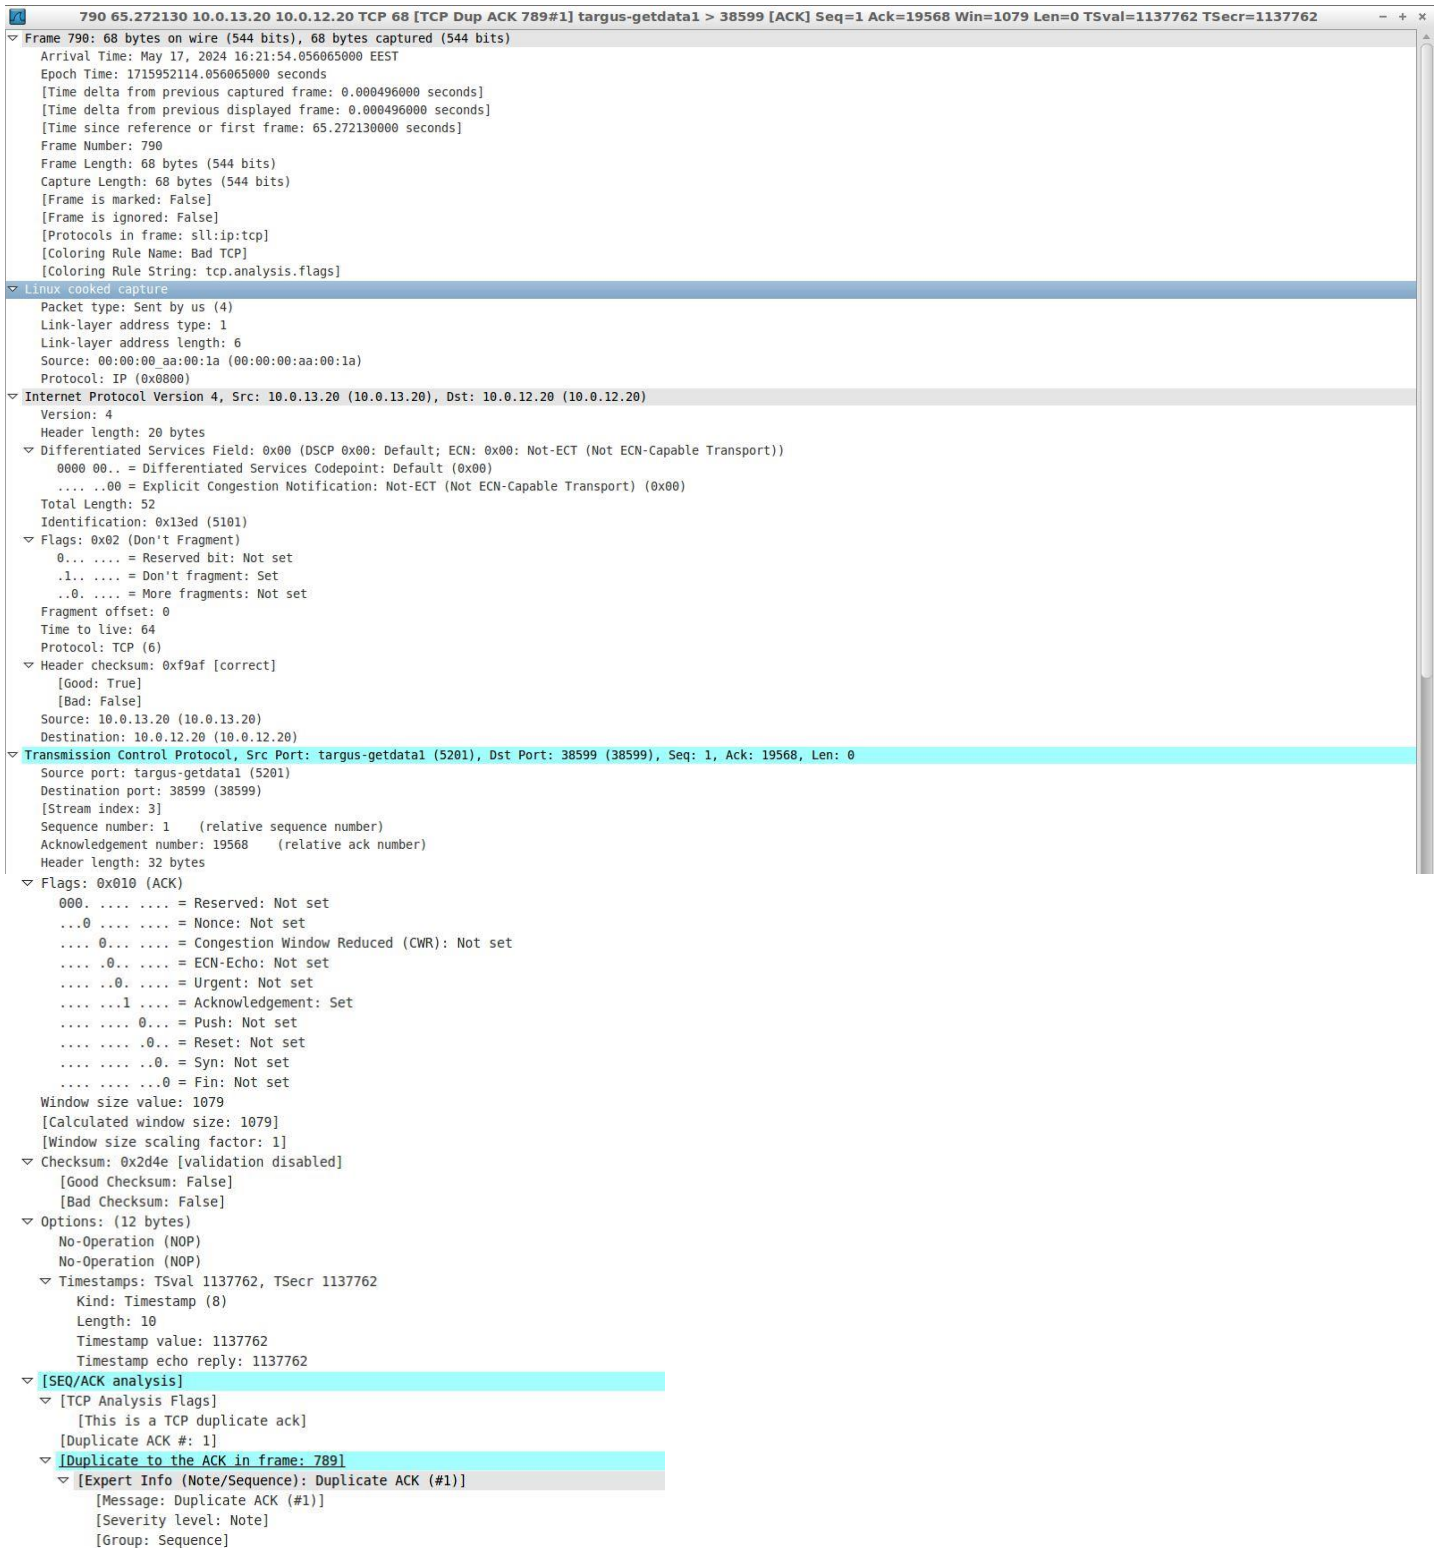


Figure 14: ACK Packet

### Comment on Your findings:

- Frame Length: → 68 bytes.
- IP Header: → 20 bytes
- TCP Header: → 32 bytes (12 Option + 20 Without Option)
- TCP Options: → 12 bytes (2 NOP + 10 Timestamp)
- TCP Options Timestamp: → 10 bytes (4T<sub>x</sub> + 4R<sub>x</sub> + 1 Length + 1 Kind)
- Ethernet Header: = Frame Length – IP Header – TCP Header = 68 – 20 – 32 = 16 bytes.
- As Shown in Fig. 14: Packet Sent from n11 to n7 has Small Size (68 B) as it contains Headers only.

## 1.2. TCP short VS long paths

- Run an iperf3 server on node n8.
- Start an iperf3 client on node n7 connecting to server on node n8 for a duration of 40 seconds and reporting interval of 10 with window size 4K.

**1.2.1. Compare the result of throughput with the case when connection was made to node n11. Why throughput drops when connecting to n8 although capacities on the two paths are the same.**

<pre>root@n7:/tmp/pycore,34478/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38566 connected to 10.0.13.20 port 5201 [ ID] Interval      Transfer    Bandwidth  Retr  Cwnd [ 4]  0.00-10.00  sec  6.93 MBytes  5.82 Mbits/sec  0    14.1 KBytes [ 4] 10.00-20.00  sec  6.93 MBytes  5.81 Mbits/sec  0    14.1 KBytes [ 4] 20.00-30.00  sec  6.92 MBytes  5.80 Mbits/sec  0    14.1 KBytes [ 4] 30.00-40.00  sec  6.94 MBytes  5.82 Mbits/sec  0    14.1 KBytes ----- [ ID] Interval      Transfer    Bandwidth  Retr [ 4]  0.00-40.00  sec  27.7 MBytes  5.81 Mbits/sec  0 [ 4]  0.00-40.00  sec  27.7 MBytes  5.81 Mbits/sec iperf Done. root@n7:/tmp/pycore,34478/n7.conf#</pre>									
<pre>root@n7:/tmp/pycore,34514/n7.conf# iperf3 -c 10.0.11.20 -t 40 -i 10 -w 4k Connecting to host 10.0.11.20, port 5201 [ 4] local 10.0.12.20 port 50864 connected to 10.0.11.20 port 5201 [ ID] Interval      Transfer    Bandwidth  Retr  Cwnd [ 4]  0.00-10.00  sec  4.53 MBytes  3.80 Mbits/sec  0    14.1 KBytes [ 4] 10.00-20.00  sec  4.49 MBytes  3.77 Mbits/sec  0    14.1 KBytes [ 4] 20.00-30.00  sec  4.43 MBytes  3.72 Mbits/sec  0    14.1 KBytes [ 4] 30.00-40.00  sec  4.57 MBytes  3.83 Mbits/sec  0    14.1 KBytes ----- [ ID] Interval      Transfer    Bandwidth  Retr [ 4]  0.00-40.00  sec  18.0 MBytes  3.78 Mbits/sec  0 [ 4]  0.00-40.00  sec  18.0 MBytes  3.78 Mbits/sec iperf Done. root@n7:/tmp/pycore,34514/n7.conf#</pre>									

Figure 15: Short Path (n7 to n11) Throughput = 5.81 Mbps

Figure 16: Long Path (n7 to n8) Throughput = 3.78 Mbps

**Comment in your findings here:**

- Throughput from n7 to n11 (Short Path): = 5.81 Mbps
- Throughput from n7 to n8 (Long Path): = 3.78 Mbps
- We observe that when n8 starts serving instead of n11, the throughput decreases because the path from n7 to n8 is longer than the path from n7 to n11, resulting in delay increase (RTT and Time out is Increased).
- By increasing Time out with the same window size, the network utilization is decreased. So, throughput decreases as shown in the Fig. 15 & Fig. 16.

## 1.3. Higher link Capacity with Drops VS Reliable Lower Capacity

This questions for this part are based on the path between n7 and n11.

For each case of the following, run iperf3 client from n7 to n11 with window size 4K.

- A Select link between n4 and n5, configure it to have capacity of 10 Mbps with zero loss in both directions.
- B Select link between n4 and n5, configure it to have capacity of 3 Mbps with zero loss in both directions.
- C Select link between n4 and n5, configure it to have capacity of 10 Mbps with 5% loss in both directions.
- D Select link between n4 and n5, configure it to have capacity of 100 Mbps with 10% loss in both directions.
- E Select link between n4 and n5, configure it to have capacity of 10 Mbps with 1% loss in direction from n4 to n5 and 0% loss in the other direction.
- F Select link between n4 and n5, configure it to have capacity of 10 Mbps with 0% loss in direction from n4 to n5 and 1% loss in the other direction.

Case	Throughput	Figures
A	5.81 Mbps	<pre> root@n7:/tmp/pycore.34478/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38566 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec     6.93 MBytes  5.82 Mbits/sec    0   14.1 KBytes [ 4] 10.00-20.00  sec     6.93 MBytes  5.81 Mbits/sec    0   14.1 KBytes [ 4] 20.00-30.00  sec     6.92 MBytes  5.80 Mbits/sec    0   14.1 KBytes [ 4] 30.00-40.00  sec     6.94 MBytes  5.82 Mbits/sec    0   14.1 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec    27.7 MBytes  5.81 Mbits/sec    0 [ 4]  0.00-40.00  sec    27.7 MBytes  5.81 Mbits/sec    0 sender receiver  iperf Done. root@n7:/tmp/pycore.34478/n7.conf# </pre>
B	2.85 Mbps	<pre> root@n7:/tmp/pycore.34523/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38611 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec     3.39 MBytes  2.85 Mbits/sec    0   14.1 KBytes [ 4] 10.00-20.00  sec     3.41 MBytes  2.86 Mbits/sec    0   14.1 KBytes [ 4] 20.00-30.00  sec     3.33 MBytes  2.84 Mbits/sec    0   14.1 KBytes [ 4] 30.00-40.00  sec     3.41 MBytes  2.86 Mbits/sec    0   14.1 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec    13.6 MBytes  2.85 Mbits/sec    0 [ 4]  0.00-40.00  sec    13.6 MBytes  2.85 Mbits/sec    0 sender receiver  iperf Done. root@n7:/tmp/pycore.34523/n7.conf# </pre>
C	493 kbps	<pre> root@n7:/tmp/pycore.34526/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38614 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec       737 KBytes  603 Kbits/sec   39   4.24 KBytes [ 4] 10.00-20.00  sec     663 KBytes  548 Kbits/sec   38   4.24 KBytes [ 4] 20.00-30.00  sec     696 KBytes  570 Kbits/sec   38   4.24 KBytes [ 4] 30.00-40.00  sec     305 KBytes  250 Kbits/sec   22   2.83 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec    2.35 MBytes  493 Kbits/sec  137 [ 4]  0.00-40.00  sec    2.34 MBytes  492 Kbits/sec  137 sender receiver  iperf Done. root@n7:/tmp/pycore.34526/n7.conf# </pre>
D	132 kbps	<pre> root@n7:/tmp/pycore.34530/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38618 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec       253 KBytes  131 Kbits/sec   25   4.24 KBytes [ 4] 10.00-20.00  sec     33.9 KBytes  27.8 Kbits/sec    9   4.24 KBytes [ 4] 20.00-30.00  sec       120 KBytes  98.5 Kbits/sec   10   4.24 KBytes [ 4] 30.00-40.00  sec       263 KBytes  215 Kbits/sec   25   4.24 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec     650 KBytes  133 Kbits/sec   69 [ 4]  0.00-40.00  sec     645 KBytes  132 Kbits/sec   69 sender receiver  iperf Done. root@n7:/tmp/pycore.34530/n7.conf# </pre>
E	3.56 Mbps	<pre> root@n7:/tmp/pycore.34533/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38621 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec     4.13 MBytes  3.47 Mbits/sec   18   8.48 KBytes [ 4] 10.00-20.00  sec     4.72 MBytes  3.96 Mbits/sec   15   8.48 KBytes [ 4] 20.00-30.00  sec     3.70 MBytes  3.10 Mbits/sec   20   8.48 KBytes [ 4] 30.00-40.00  sec     4.42 MBytes  3.71 Mbits/sec   17   8.48 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec    17.0 MBytes  3.56 Mbits/sec   70 [ 4]  0.00-40.00  sec    17.0 MBytes  3.56 Mbits/sec   70 sender receiver  iperf Done. root@n7:/tmp/pycore.34533/n7.conf# </pre>
F	2.8 Mbps	<pre> root@n7:/tmp/pycore.34537/n7.conf# iperf3 -c 10.0.13.20 -t 40 -i 10 -w 4k Connecting to host 10.0.13.20, port 5201 [ 4] local 10.0.12.20 port 38625 connected to 10.0.13.20 port 5201 [ ID] Interval           Transfer     Bandwidth   Retr  Cwnd [ 4]  0.00-10.00  sec     3.82 MBytes  3.21 Mbits/sec   22   5.66 KBytes [ 4] 10.00-20.00  sec     3.34 MBytes  2.80 Mbits/sec   25   8.48 KBytes [ 4] 20.00-30.00  sec     2.78 MBytes  2.33 Mbits/sec   29   4.24 KBytes [ 4] 30.00-40.00  sec     3.41 MBytes  2.86 Mbits/sec   24   4.24 KBytes ----- [ ID] Interval           Transfer     Bandwidth   Retr [ 4]  0.00-40.00  sec    13.3 MBytes  2.80 Mbits/sec  100 [ 4]  0.00-40.00  sec    13.3 MBytes  2.80 Mbits/sec  100 sender receiver  iperf Done. root@n7:/tmp/pycore.34537/n7.conf# </pre>



### 1.3.1. Compare throughputs in cases a, b, c. Why b is better than c?

Comment in your findings here:

We observe that the throughput in case (A) is higher than in case (B), and the throughput in case (B) is higher than in case (C). This indicates that case (B) performs better than case (C). Although the link capacity in case (C) '10 Mbps' is greater than in case (B) '3Mbps' but in case (C) there are higher losses (5%) that necessitate retransmissions, unlike in case (B) that has loss rate is 0%. In addition to that, the throughput in case (A) is greater than that of case (B) due to the greater link capacity in case (A) compared to case (B) as both case (A) and case (B) have equal losses.

### 1.3.2. Compare throughputs in cases b, c and d. Which is better? Why?

Comment in your findings here:

Case (B) is better than cases (C) and (D) [ $\text{Throughput}_B > \text{Throughput}_C > \text{Throughput}_D$ ] because the loss in case (B) is zero in both directions, eliminating the need for the client or server to retransmit any packets. Although the capacity in case (B) is lower than in cases (C) and (D), packet retransmissions are necessary in cases (C) and (D) due to non-zero loss (for case (C) = 5% loss & for case (D) = 10% loss).

### 1.3.3. Compare throughputs in cases e and f? Which is better? Why?

Comment in your findings here:

Throughput for Case (E) is better than case (F) because large packets are often sent from n5 to n4, while the packets sent from n4 to n5 are usually smaller. Consequently, losses will affect the (n5 to n4) link more than (n4 to n5) link.

## 2. OSPF Lab

Each of the following sections describes an experiment based on the network topology provided by the [Project OSPF TCP.imn](#) file.

The [Introductory Lab](#) explained and introduced most of the needed tools like iperf3, vtysh, wireshark, etc. If you have not gone through it, do not start working on the assignment below.

It is better to start the network emulation once and run Wireshark captures. You can reset Wireshark captures from experiment to another.

### 2.1. OSPF Link Cost Change

1. Stop any running iperf3 clients.
2. Set all links to have zero loss in the two directions with 100 Mbps speed.
3. Run iperf3 between n7 and n11 for a duration of 500 seconds or longer. Identify the path between n7 and n11.
4. Open vtysh on node n5 by opening a bash terminal and typing vtysh. Now we can configure the router and link costs.
5. Type the following in vtysh: show Ip ospf interface eth1 (To know the ethernet number of each interface you can choose toolbar->view->show->interface names) This displays information about interface eth1. Note its ospf cost.
6. Type the following configure terminal.
  - interface eth1 (the interface for the link between n5 and n4)
  - ospf cost 40 (set cost to 40)

**2.1.1. Check what happens to the path between n7 and n11 (as seen after steps 3 and 6)? Explain what happens. To help visualize the path choose toolbar->widgets-> Throughput.**

```
root@n5:/tmp/pycore.41071/n5.conf# vtysh
Hello, this is Quagga (version 0.99.20.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

n5# show ip ospf interface eth1
eth1 is up
ifindex 116, MTU 1500 bytes, BW 0 Kbit <UP,BROADCAST,RUNNING,MULTICAST>
Internet Address 10.0.4.2/24, Area 0.0.0.0
MTU mismatch detection:enabled
Router ID 10.0.5.1, Network Type BROADCAST, Cost: 10
Transmit Delay is 1 sec, State DR, Priority 1
Designated Router (ID) 10.0.5.1, Interface Address 10.0.4.2
Backup Designated Router (ID) 10.0.3.1, Interface Address 10.0.4.1
Multicast group memberships: OSPFAllRouters OSPFDesignatedRouters
Timer intervals configured, Hello 10s, Dead 40s, Wait 40s, Retransmit 5
Hello due in 7.095s
Neighbor Count is 1, Adjacent neighbor count is 1
n5#
```

Figure 17: n5 interface 1




Figure 18: Path From n7 to n11 before changing n5 cost.




Figure 19: Path From n7 to n11 after changing n5 cost to 40

**Comment on your findings:**

As illustrated in Fig. 18 & 19, increasing the cost of the interface between n4 and n5 causes the data packets to take an alternative route to achieve a lower overall path cost, with the help of Dijkstra's Algorithm.

The routers begin to recalculate the best route to determine the shortest path between n7 and n11 as shown in Fig. 19

**2.1.2. Set the cost of eth1 at node n5 back to 10. Establish two iperf3 connections: one from n7 to n11 and the second from n11 to n7 both for duration of 500 seconds. Now go to node n4 and set interface cost for interface connecting n4 with n5 to 40. What happens in the paths of the two connections? Explain what happens. What do you conclude?**

```

root@n4:/tmp/pycore.59083/n4.conf# vtysh

Hello, this is Quagga (version 0.99.20.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

n4# show ip ospf interface eth1
eth1 is up
  ifindex 361, MTU 1500 bytes, BW 0 Kbit <UP,BROADCAST,RUNNING,MULTICAST>
  Internet Address 10.0.4.1/24, Area 0.0.0.0
  MTU mismatch detection:enabled
  Router ID 10.0.3.1, Network Type BROADCAST, Cost: 10
  Transmit Delay is 1 sec, State Backup, Priority 1
  Designated Router (ID) 10.0.5.1, Interface Address 10.0.4.2
  Backup Designated Router (ID) 10.0.3.1, Interface Address 10.0.4.1
  Multicast group memberships: OSPFAllRouters OSPFDesignatedRouters
  Timer intervals configured, Hello 10s, Dead 40s, Wait 40s, Retransmit 5
  Hello due in 3.903s
  Neighbor Count is 1, Adjacent neighbor count is 1
n4#

```

Figure 20: n4 interface 1




Figure 21: Path from n7 to n11 and vice versa before changing n4 cost.




Figure 22: Path from n7 to n11 and vice versa after changing n4 cost to 40

### Comment on your findings:

For the same reason mentioned above, the path becomes much longer. The increasing cost of a link forces the router to choose an alternative path, even if it is longer, to reduce congestion. As OSPF could handle such situations by increasing the cost of  $n4 \rightarrow n5$  and leaving  $n5 \rightarrow n4$  the same. The shortest path from n7 to n11 is not the same as the shortest path from n11 to n7. Paths are  $(n7 \rightarrow n5 \rightarrow n4 \rightarrow n9 \rightarrow n11)$  &  $(n11 \rightarrow n9 \rightarrow n10 \rightarrow n6 \rightarrow n5 \rightarrow n7)$



## 2.2. OSPF Database Updates

1. Start Wireshark and have its filter to capture only OSPF related packets.
2. Go to another router say router n2, open vtysh and issue the commands:
  - config terminal
  - show ip ospf database ([check this link to understand more](#))
  - show ip ospf route.
3. On router n4, set link cost of eth1 to 20. Capture the link state packets advertised in Wireshark.
4. Go to router n4, go out of vtysh, or open new bash terminal and issue the following commands to disconnect router n4 from the network:
  - ifconfig eth0 down
  - ifconfig eth1 down
  - ifconfig eth2 down

### 2.2.1. Capture and explain the outputs due to execution of step 2. Why some destinations have more than one route in the routing tables?

```
Hello, this is Quagga (version 0.99.20.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

n2# show ip ospf database

      OSPF Router with ID (10.0.1.2)

      Router Link States (Area 0.0.0.0)

Link ID      ADV Router   Age  Seq#       CkSum  Link count
10.0.0.1     10.0.0.1     623  0x8000000a 0xf2bd 3
10.0.0.2     10.0.0.2     623  0x8000000b 0x9809 3
10.0.1.2     10.0.1.2     623  0x8000000a 0x8d0a 3
10.0.3.1     10.0.3.1     615  0x80000009 0xc1ce 3
10.0.3.2     10.0.3.2     616  0x8000000a 0xadde 3
10.0.5.1     10.0.5.1     621  0x80000008 0x9405 3
10.0.6.2     10.0.6.2     620  0x80000009 0xccab 3
10.0.8.2     10.0.8.2     619  0x8000000b 0x2414 4

      Net Link States (Area 0.0.0.0)

Link ID      ADV Router   Age  Seq#       CkSum
10.0.0.2     10.0.0.2     624  0x80000001 0x5fcb
10.0.1.2     10.0.1.2     623  0x80000001 0x64c2
10.0.2.1     10.0.1.2     623  0x80000001 0x51d6
10.0.3.2     10.0.3.2     616  0x80000001 0x6bb3
10.0.4.2     10.0.5.1     621  0x80000001 0x64b7
...skipping one line
10.0.6.2     10.0.6.2     625  0x80000001 0x4bcc
10.0.7.1     10.0.6.2     625  0x80000001 0x59bb
10.0.8.2     10.0.8.2     619  0x80000001 0x52bd
10.0.9.1     10.0.8.2     624  0x80000001 0x2de3
10.0.10.2    10.0.8.2     619  0x80000001 0x61a8
```

Figure 23: n2 Database

```
n2# show ip ospf route
===== OSPF network routing table =====
N  10.0.0.0/24      [20] area: 0.0.0.0
                        via 10.0.1.1, eth0
                        via 10.0.2.2, eth1
N  10.0.1.0/24      [10] area: 0.0.0.0
                        directly attached to eth0
N  10.0.2.0/24      [10] area: 0.0.0.0
                        directly attached to eth1
N  10.0.3.0/24      [30] area: 0.0.0.0
                        via 10.0.3.1, eth2
N  10.0.4.0/24      [30] area: 0.0.0.0
                        via 10.0.3.1, eth2
N  10.0.5.0/24      [40] area: 0.0.0.0
                        via 10.0.1.1, eth0
                        via 10.0.3.1, eth2
N  10.0.6.0/24      [20] area: 0.0.0.0
                        via 10.0.1.1, eth0
N  10.0.7.0/24      [30] area: 0.0.0.0
                        via 10.0.1.1, eth0
                        via 10.0.3.1, eth2
N  10.0.8.0/24      [20] area: 0.0.0.0
                        via 10.0.3.1, eth2
N  10.0.9.0/24      [10] area: 0.0.0.0
                        directly attached to eth2
N  10.0.10.0/24     [20] area: 0.0.0.0
                        via 10.0.3.1, eth2
N  10.0.11.0/24     [20] area: 0.0.0.0
                        via 10.0.2.2, eth1
N  10.0.12.0/24     [40] area: 0.0.0.0
                        via 10.0.3.1, eth2
N  10.0.13.0/24     [20] area: 0.0.0.0
                        via 10.0.3.1, eth2

===== OSPF router routing table =====
===== OSPF external routing table =====
```

Figure 24: n2 Routing table.

#### Comment on your findings:

From the above routing table there are multiple paths available that have the same cost. This means that the network switch can choose any one of these paths to send data. Additionally, if one of these paths becomes unavailable (goes down), there are still other paths that can be used, ensuring continued network connectivity.

The database for router n2 is illustrated in the figures, and it consists of two segments:

Router (n2) Link State: This describes the network's routers, links, and topology, constructed at each router after exchanging hello packets with other routers.

Designated Router Link State: Similar to the router link state but generated by the designated router, which is elected to manage communication on a multi-access network segment, reducing overhead and ensuring efficient data exchange.

Figures 23 and 24 show the n2 database and routing table, respectively. Both tables have fields including:

Link ID: Identifies the router or network for which the LSA is originated.

ADV Router: Specifies the originating router of the LSA.

Age: Indicates the time since the LSA was last updated.

SEQ#: The sequence number of the LSA, used to compare and accept new information.



Checksum: Validates the LSA packet.

Link Count: Number of links advertised within the LSA.

The routing table in figure 24 includes destination addresses and their corresponding routes. OSPF supports Equal Cost Multipath (ECMP), meaning multiple paths can be assigned to the same destination if they have equal cost, balancing load and increasing network fault tolerance. Some destinations have multiple equal-cost paths, while others have a single path.

### 2.2.2. After executing step 3, determine how long it took the network to exchange link state packets and adjust routing tables. (Hint: you can calculate the required time by observing the time of first OSPF update message and the last ACK from Wireshark).

No.	Time	Source	Destination	Protocol	Length	Info
10322	1380.229640	10.0.0.2	224.0.0.5	OSPF	84	Hello Packet
10323	1380.229645	10.0.0.1	224.0.0.5	OSPF	84	Hello Packet
10324	1380.229647	10.0.0.1	224.0.0.5	OSPF	84	Hello Packet
10325	1388.370918	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10326	1388.370924	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10327	1388.370933	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10328	1388.370936	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10329	1388.371087	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10330	1388.371090	10.0.0.1	224.0.0.5	OSPF	208	LS Update
10331	1388.371166	10.0.0.1	224.0.0.5	OSPF	112	LS Update
10332	1388.371170	10.0.0.1	224.0.0.5	OSPF	112	LS Update
10333	1388.371179	10.0.0.2	224.0.0.5	OSPF	112	LS Update
10334	1388.371182	10.0.0.2	224.0.0.5	OSPF	112	LS Update
10335	1388.371224	10.0.0.1	224.0.0.5	OSPF	112	LS Update
10336	1388.371226	10.0.0.1	224.0.0.5	OSPF	112	LS Update
10337	1388.371335	10.0.0.1	224.0.0.5	OSPF	112	LS Update

Figure 25: n2 First Update Packet

No.	Time	Source	Destination	Protocol	Length	Info
10389	1388.730011	10.0.0.2	224.0.0.5	OSPF	80	LS Acknowledge
10390	1388.730014	10.0.0.2	224.0.0.5	OSPF	80	LS Acknowledge
10391	1388.730439	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
10392	1388.730443	10.0.0.2	224.0.0.5	OSPF	80	LS Acknowledge
10393	1388.884656	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
10394	1388.884671	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
10395	1388.885069	10.0.0.1	224.0.0.5	OSPF	80	LS Acknowledge
10396	1390.108222	10.0.0.2	224.0.0.5	OSPF	84	Hello Packet
10397	1390.108222	10.0.0.1	224.0.0.5	OSPF	84	Hello Packet
10398	1390.108231	10.0.0.2	224.0.0.5	OSPF	84	Hello Packet
10399	1390.108237	10.0.0.1	224.0.0.5	OSPF	84	Hello Packet
10400	1390.108246	10.0.0.2	224.0.0.5	OSPF	84	Hello Packet
10401	1390.108252	10.0.0.1	224.0.0.5	OSPF	84	Hello Packet
10402	1390.108253	10.0.0.2	224.0.0.5	OSPF	84	Hello Packet

Figure 26: n2 Last ACK Packet

#### Comment on your findings:

From the above graphs we conclude that the time taken by the Network to exchange link state packets and adjust routing tables → Time taken = 1388.885069 – 1388.370918 = 0.514151 sec

### 2.2.3. After execution of step 4, identify the new routing table and router database at router n2. Explain the updates in the new routing table and the new database.

```
Hello, this is Quagga (version 0.99.20.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

n2# show ip ospf database

    OSPF Router with ID (10.0.1.2)

    Router Link States (Area 0.0.0.0)

Link ID      ADV Router   Age Seq#       CkSum Link count
10.0.0.1     10.0.0.1     1307 0x8000000b 0xf0be 3
10.0.0.2     10.0.0.2     1357 0x8000000c 0x960a 3
10.0.1.2     10.0.1.2     1356 0x8000000b 0x8b0b 3
10.0.3.1     10.0.3.1     1311 0x8000000e 0xba3f 1
10.0.3.2     10.0.3.2     1669 0x8000000b 0xf4a6 3
10.0.5.1     10.0.5.1     1280 0x8000000a 0xffa8 3
10.0.6.2     10.0.6.2     1288 0x8000000a 0xcaac 3
10.0.8.2     10.0.8.2     1267 0x8000000d 0x2823 4

    Net Link States (Area 0.0.0.0)

Link ID      ADV Router   Age Seq#       CkSum
10.0.0.2     10.0.0.2     737 0x80000002 0x5dcc
10.0.1.2     10.0.1.2     736 0x80000002 0x62c3
10.0.2.1     10.0.1.2     1656 0x80000002 0x4fd7
10.0.5.1     10.0.5.1     720 0x80000002 0x67b2
10.0.6.2     10.0.6.2     78 0x80000003 0x47ce
...skipping one line
10.0.9.1     10.0.8.2     707 0x80000002 0x2be4
10.0.10.2    10.0.8.2     257 0x80000003 0x5daa
```

Figure 27: n2 Database after Shutting down n4

```
n2# show ip ospf route
===== OSPF network routing table =====
N   10.0.0.0/24      [20] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.2.2, eth1
N   10.0.1.0/24      [10] area: 0.0.0.0
    directly attached to eth0
N   10.0.2.0/24      [10] area: 0.0.0.0
    directly attached to eth1
N   10.0.3.0/24      [40] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.4.0/24      [50] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.5.0/24      [40] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.6.0/24      [20] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.7.0/24      [30] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.8.0/24      [20] area: 0.0.0.0
    via 10.0.9.1, eth2
N   10.0.9.0/24      [10] area: 0.0.0.0
    directly attached to eth2
N   10.0.10.0/24     [20] area: 0.0.0.0
    via 10.0.9.1, eth2
N   10.0.11.0/24     [20] area: 0.0.0.0
    via 10.0.2.2, eth1
N   10.0.12.0/24     [50] area: 0.0.0.0
    via 10.0.1.1, eth0
    via 10.0.9.1, eth2
N   10.0.13.0/24     [20] area: 0.0.0.0
    via 10.0.9.1, eth2

===== OSPF router routing table =====
===== OSPF external routing table =====
```

Figure 28: n2 Routing Table after Shutting down n4

#### Comment on your findings:

If the node (n4) is disconnected, another node will take its place which is (n2). Although n2 has a higher cost than n4, the network will update its database to reflect this change. This ensures that data can still be routed effectively, even though the overall path cost will be higher.